

# buy\_affinity

July 26, 2020

## 1 Case Study - Customer Buy Affinity On E-commerce Dataset

```
[2]: #importing for exploratory data analysis and visualization
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

#importing for undersampling and oversampling techniques
import numpy as np
from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import RandomOverSampler

from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from sklearn.preprocessing import StandardScaler

#different models to try
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
```

## 2 Reading the data for training

```
[3]: df = pd.read_csv('BuyAffinity_Train.txt', sep='\t', index_col='Index')
```

## 3 Checking the head of data

```
[4]: df.head()
```

```
[4]:
```

	F1	F2	F3	F4	F5	F6	F7	F8	F9	\
Index										
1	0.224506	0.500340	0.489860	0.902413	7934	-6970	-5714	9982	-5697	
2	0.321128	0.281119	0.907283	0.772159	-8238	1219	1663	1287	-3658	
3	0.893441	0.622005	0.998776	0.098386	8540	5266	-9377	-3504	-4511	
4	0.320641	0.957234	0.346000	0.646479	-7772	-383	9681	-8661	3474	

```
5      0.475961  0.623008  0.544988  0.159709  1571 -8039 -7961 -2385  4407
```

```

      F10  ...      F14      F15      F16  F17 F18  F19  F20  \
Index  ...
1      4227810299  ... -3433637453  10/4/1986    9/6/1992    2   1  706  305
2     -1146724819  ...   609277486  2/24/1979    1/5/1983    1   1  423  206
3      5947184989  ... -8977995005  1/12/1989   11/22/1986    2   1  703  315
4     -5724795826  ...  4868760308  2/18/1982    6/10/1992    1   1  122  304
5     -3097637172  ...  9757408267  4/10/1987   10/19/1985    1   1  486  240

```

```

      F21  F22  C
Index
1         1    2  0
2        18    7  1
3         1    4  0
4        15    1  0
5         1    1  0

```

[5 rows x 23 columns]

## 4 Dropping the F14,F15 date columns

```
[5]: df = df.drop(columns=['F15', 'F16'])
      df.head()
```

```
[5]:
      F1      F2      F3      F4      F5      F6      F7      F8      F9  \
Index
1      0.224506  0.500340  0.489860  0.902413  7934 -6970 -5714  9982 -5697
2      0.321128  0.281119  0.907283  0.772159 -8238  1219  1663  1287 -3658
3      0.893441  0.622005  0.998776  0.098386  8540  5266 -9377 -3504 -4511
4      0.320641  0.957234  0.346000  0.646479 -7772  -383  9681 -8661  3474
5      0.475961  0.623008  0.544988  0.159709  1571 -8039 -7961 -2385  4407

```

```

      F10  ...      F12      F13      F14  F17  F18  F19  \
Index  ...
1      4227810299  ...   316195953  6176861823 -3433637453    2   1  706
2     -1146724819  ...  1378635942 -9031507610   609277486    1   1  423
3      5947184989  ... -9921889287 -5610051842 -8977995005    2   1  703
4     -5724795826  ...  6550322883 -4697085930  4868760308    1   1  122
5     -3097637172  ...   759031103  9984692447  9757408267    1   1  486

```

```

      F20  F21  F22  C
Index
1      305    1    2  0
2      206   18    7  1
3      315    1    4  0

```

```

4      304   15    1   0
5      240    1    1   0

```

[5 rows x 21 columns]

## 5 Describing the data

```
[6]: df.describe()
```

```

[6]:
      count      F1      F2      F3      F4  \
count  101180.000000  101180.000000  101180.000000  101180.000000
mean      0.502348      0.501497      0.499886      0.499839
std      0.288058      0.289017      0.288875      0.288729
min      0.000018      0.000004      0.000002      0.000006
25%      0.253819      0.251115      0.248818      0.250501
50%      0.501802      0.501095      0.499820      0.501387
75%      0.753598      0.752404      0.750281      0.748803
max      0.999986      0.999990      0.999985      0.999977

      count      F5      F6      F7      F8  \
count  101180.000000  101180.000000  101180.000000  101180.000000
mean     -29.742617      1.511000      7.939118     16.434147
std     5781.829379     5796.594007     5765.581875     5782.805211
min    -10000.000000    -10000.000000    -10000.000000    -10000.000000
25%    -5045.000000    -5012.000000    -4979.000000    -4988.000000
50%     -46.000000     -11.500000      16.500000      60.000000
75%     4978.000000     5050.000000     4962.000000     5021.000000
max     10000.000000     10000.000000     10000.000000     10000.000000

      count      F9      F10  ...      F12      F13  \
count  101180.000000  1.011800e+05  ...  1.011800e+05  1.011800e+05
mean      9.927812  1.869749e+07  ...  6.921685e+06  1.127538e+07
std      5771.004738  5.769064e+09  ...  5.771750e+09  5.786659e+09
min    -10000.000000  -9.999816e+09  ...  -9.999754e+09  -9.999852e+09
25%    -4992.250000  -4.966961e+09  ...  -4.991278e+09  -5.027003e+09
50%      22.500000  4.346691e+07  ...  4.277384e+07  9.416236e+06
75%     5020.000000  5.019440e+09  ...  5.008238e+09  5.032947e+09
max     10000.000000  9.999497e+09  ...  9.999555e+09  9.999951e+09

      count      F14      F17      F18      F19  \
count  1.011800e+05  101180.000000  101180.000000  101180.000000
mean   -5.198163e+07      1.339593      1.335027      387.181479
std     5.782358e+09      0.894065      0.882435      235.939585
min    -9.999767e+09      1.000000      1.000000      1.000000
25%    -5.070403e+09      1.000000      1.000000      187.000000
50%    -8.816174e+07      1.000000      1.000000      374.000000

```

75%	4.947133e+09	1.000000	1.000000	560.000000
max	9.999997e+09	6.000000	6.000000	901.000000

	F20	F21	F22	C
count	101180.000000	101180.000000	101180.000000	101180.000000
mean	387.333999	4.588031	4.598567	0.245375
std	235.374285	4.775671	4.776494	0.430311
min	1.000000	1.000000	1.000000	0.000000
25%	187.000000	1.000000	1.000000	0.000000
50%	375.000000	3.000000	3.000000	0.000000
75%	561.000000	7.000000	7.000000	0.000000
max	901.000000	21.000000	21.000000	1.000000

[8 rows x 21 columns]

## 6 Finding the percentage missing data

No data is missing for any attribute


```
[7]: miss_per = df.isnull().sum()/df.shape[0]
      print(miss_per)
```

```
F1      0.0
F2      0.0
F3      0.0
F4      0.0
F5      0.0
F6      0.0
F7      0.0
F8      0.0
F9      0.0
F10     0.0
F11     0.0
F12     0.0
F13     0.0
F14     0.0
F17     0.0
F18     0.0
F19     0.0
F20     0.0
F21     0.0
F22     0.0
C       0.0
dtype: float64
```

## 7 Checking distribution of positive and negative class

There is class imbalance between negative and positive class

```
[8]: _=sns.countplot(x='C',data=df)
```



## 8 Training with LogisticRegression

Does not work well on imbalanced data

```
[9]: data = df.to_numpy()
X,y = data[:, :-1], data[:, -1]
X =StandardScaler().fit_transform(X)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
↪30, stratify=y)
clf = LogisticRegression(verbose=1)
clf.fit(X_train,y_train)
y_pred = clf.predict(X_test)
rep=classification_report(y_test, y_pred)
print(rep)
```

precision	recall	f1-score	support
-----------	--------	----------	---------

	0.0	0.75	1.00	0.86	22906
	1.0	0.00	0.00	0.00	7448
accuracy				0.75	30354
macro avg	0.38	0.50	0.43		30354
weighted avg	0.57	0.75	0.65		30354

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.1s finished
/home/durgesh/miniconda3/lib/python3.7/site-
packages/sklearn/metrics/_classification.py:1221: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
```

## 9 Training With RandomForest

Performs a bit better to logistic regression on imbalanced data

```
[10]: data = df.to_numpy()
X,y = data[:, :-1], data[:, -1]
X = StandardScaler().fit_transform(X)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
↪30, stratify=y)
clf = RandomForestClassifier(verbose=1, n_jobs=80)
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
rep = classification_report(y_test, y_pred)
print(rep)
```

```
[Parallel(n_jobs=80)]: Using backend ThreadingBackend with 80 concurrent
workers.
[Parallel(n_jobs=80)]: Done 42 out of 100 | elapsed: 14.4s remaining: 19.9s
[Parallel(n_jobs=80)]: Done 100 out of 100 | elapsed: 18.2s finished
[Parallel(n_jobs=80)]: Using backend ThreadingBackend with 80 concurrent
workers.
```

	precision	recall	f1-score	support	
	0.0	0.76	0.99	0.86	22906
	1.0	0.44	0.02	0.04	7448
accuracy				0.75	30354
macro avg	0.60	0.51	0.45		30354
weighted avg	0.68	0.75	0.66		30354

```
[Parallel(n_jobs=80)]: Done 42 out of 100 | elapsed: 0.4s remaining: 0.6s
```

```
[Parallel(n_jobs=80)]: Done 100 out of 100 | elapsed:    0.4s finished
```

## 10 Performing Under Sampling And Training With RandomForest

Random forests performs better with balanced data

```
[11]: rus = RandomUnderSampler()
data = df.to_numpy()
X,y = data[:, :-1], data[:, -1]
X = StandardScaler().fit_transform(X)
X,y = rus.fit_resample(X, y)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
↪ 30, stratify=y)
clf = RandomForestClassifier(verbose=1, n_jobs=80)
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
rep = classification_report(y_test, y_pred)
print(rep)
```

```
[Parallel(n_jobs=80)]: Using backend ThreadingBackend with 80 concurrent workers.
```

```
[Parallel(n_jobs=80)]: Done 42 out of 100 | elapsed:    6.4s remaining:    8.8s
```

```
[Parallel(n_jobs=80)]: Done 100 out of 100 | elapsed:    8.0s finished
```

```
[Parallel(n_jobs=80)]: Using backend ThreadingBackend with 80 concurrent workers.
```

	precision	recall	f1-score	support
0.0	0.79	0.48	0.60	7449
1.0	0.63	0.88	0.73	7448
accuracy			0.68	14897
macro avg	0.71	0.68	0.66	14897
weighted avg	0.71	0.68	0.66	14897

```
[Parallel(n_jobs=80)]: Done 42 out of 100 | elapsed:    0.2s remaining:    0.3s
```

```
[Parallel(n_jobs=80)]: Done 100 out of 100 | elapsed:    0.2s finished
```

## 11 Performing Over Sampling And Training with RandomForest

Random forests performs better with over-sampling balanced data as there is more datapoints to train

```
[12]: ros = RandomOverSampler()
data = df.to_numpy()
X,y = data[:, :-1], data[:, -1]
```

```

X =StandardScaler().fit_transform(X)
X,y = ros.fit_resample(X, y)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
    ↳30,stratify=y)
clf = RandomForestClassifier(verbose=1,n_jobs=80)
clf.fit(X_train,y_train)
y_pred = clf.predict(X_test)
rep=classification_report(y_test, y_pred)
print(rep)

```

```

[Parallel(n_jobs=80)]: Using backend ThreadingBackend with 80 concurrent
workers.
[Parallel(n_jobs=80)]: Done 42 out of 100 | elapsed: 21.0s remaining: 29.0s
[Parallel(n_jobs=80)]: Done 100 out of 100 | elapsed: 26.8s finished
[Parallel(n_jobs=80)]: Using backend ThreadingBackend with 80 concurrent
workers.

```

	precision	recall	f1-score	support
0.0	0.91	0.83	0.87	22906
1.0	0.85	0.92	0.88	22906
accuracy			0.88	45812
macro avg	0.88	0.88	0.88	45812
weighted avg	0.88	0.88	0.88	45812

```

[Parallel(n_jobs=80)]: Done 42 out of 100 | elapsed: 0.7s remaining: 1.0s
[Parallel(n_jobs=80)]: Done 100 out of 100 | elapsed: 0.8s finished

```

## 12 Predicting labels on BuyAffinity\_Test data

```

[26]: df_test = pd.read_csv('BuyAffinity_Test.txt',sep='\t',index_col='Index')
df_test =df_test.drop(columns=['F15', 'F16'])
data_test = df_test.to_numpy()
X_test = data_test
X_test =StandardScaler().fit_transform(X_test)
y_pred = clf.predict(X_test).astype(np.int32)
df_test = pd.read_csv('BuyAffinity_Test.txt',sep='\t',index_col='Index')
df_test['C']= y_pred
df_test.to_csv('BuyAffinity_Test_labels.txt', sep='\t', mode='w')
df_test.head()

```

```

[Parallel(n_jobs=80)]: Using backend ThreadingBackend with 80 concurrent
workers.
[Parallel(n_jobs=80)]: Done 42 out of 100 | elapsed: 0.3s remaining: 0.5s
[Parallel(n_jobs=80)]: Done 100 out of 100 | elapsed: 0.4s finished

```



[26]:

	F1	F2	F3	F4	F5	F6	F7	F8	\
Index									
T30234341	0.654765	0.812009	0.603190	0.391039	-5220	4825	-1784	7447	
T30234342	0.694636	0.690568	0.473460	0.259760	-618	-5018	2012	9259	
T30234343	0.203759	0.323301	0.492294	0.011448	-8778	6141	6965	3774	
T30234344	0.319627	0.286247	0.906197	0.093840	-7929	4471	7715	9543	
T30234345	0.236003	0.782784	0.285689	0.383585	-3296	4564	-1580	-8559	

	F9	F10	...	F14	F15	F16	F17	F18	\
Index			...						
T30234341	-7147	-3461806391	...	5553595074	9/17/1996	8/18/1990	1	1	
T30234342	9267	-36253473	...	2216284070	11/7/1985	4/11/1990	1	1	
T30234343	4303	5354243488	...	-315409510	7/9/1984	5/4/1997	1	1	
T30234344	335	7405036171	...	-3360224957	6/3/1987	7/1/1988	1	1	
T30234345	-27	-6351599280	...	7604838279	7/17/1984	12/3/1993	1	1	

	F19	F20	F21	F22	C
Index					
T30234341	436	478	1	1	0
T30234342	138	56	10	4	0
T30234343	117	323	10	1	0
T30234344	115	149	16	21	0
T30234345	527	281	3	1	1

[5 rows x 23 columns]